

Customed Wrapper

Yehu Chen

2018/1/28

```
library(nnlS)
library(SuperLearner)

## Super Learner
## Version: 2.0-22
## Package created on 2017-07-18
# Review available models.
listWrappers()

## All prediction algorithm wrappers in SuperLearner:
## [1] "SL.bartMachine"      "SL.bayesglm"      "SL.biglasso"
## [4] "SL.caret"           "SL.caret.rpart"   "SL.cforest"
## [7] "SL.dbarts"          "SL.earth"         "SL.extraTrees"
## [10] "SL.gam"             "SL.gbm"           "SL.glm"
## [13] "SL.glm.interaction" "SL.glmnet"        "SL.ipredbagg"
## [16] "SL.kernelKnn"       "SL.knn"           "SL.ksvm"
## [19] "SL.lda"             "SL.leekasso"      "SL.lm"
## [22] "SL.loess"           "SL.logreg"        "SL.mean"
## [25] "SL.nnet"            "SL.nnlS"          "SL.polymars"
## [28] "SL.qda"             "SL.randomForest"  "SL.ranger"
## [31] "SL.ridge"           "SL.rpart"         "SL.rpartPrune"
## [34] "SL.speedglm"        "SL.speedlm"       "SL.step"
## [37] "SL.step.forward"    "SL.step.interaction" "SL.stepAIC"
## [40] "SL.svm"             "SL.template"      "SL.xgboost"

##
## All screening algorithm wrappers in SuperLearner:
## [1] "All"
## [1] "screen.corP"          "screen.corRank"    "screen.glmnet"
## [4] "screen.randomForest"  "screen.SIS"        "screen.template"
## [7] "screen.ttest"         "write.screen.template"
```

Read binomial dataset

Separate data into 50% training, 25% holdout validation and 25% test sets

```
filename = "Binomial_dataset_5_final.txt"
data = read.csv(filename)

# First divide data into positive and negative sets
data_pos = subset(data, Ytemp==1)
```

```

data_neg = subset(data, Ytemp==0)

# select the 50%, 25%, 25% to be training, holdout and testing data
data_train_pos = data_pos[c(2:as.integer(nrow(data_pos)/2)),]
data_holdout_pos = data_pos[c(as.integer(nrow(data_pos)/2+1):as.integer(3*nrow(data_pos)/4)),]
data_test_pos = data_pos[c(as.integer(3*nrow(data_pos)/4+1):nrow(data_pos)),]

data_train_neg = data_neg[c(2:as.integer(nrow(data_neg)/2)),]
data_holdout_neg = data_neg[c(as.integer(nrow(data_neg)/2+1):as.integer(3*nrow(data_neg)/4)),]
data_test_neg = data_neg[c(as.integer(3*nrow(data_neg)/4+1):nrow(data_neg)),]

# stack the pos and neg datasets
data_train = rbind(data_train_pos, data_train_neg)
data_holdout = rbind(data_holdout_pos, data_holdout_neg)
data_test = rbind(data_test_pos, data_test_neg)

# split datasets to X and Y
X_train = data_train[,c(2:ncol(data_train))]
X_test = data_test[,c(2:ncol(data_train))]
X_holdout = data_holdout[,c(2:ncol(data_train))]
Y_train = data_train[,c(1)]
Y_test = data_test[,c(1)]
Y_holdout = data_holdout[,c(1)]

```

Notice that the final training, holdout, and testing Y are still imbalanced.

We only split the dataset to training, holdout, and testing in the same proportion as the original dataset. This might require extra work, but can eliminate the pitfall where the data is imbalanced to positive but the training data is imbalanced to negative.

Customize SVM wrappers

```

library(caret)

## Loading required package: lattice
## Loading required package: ggplot2

library(lattice)
library(ggplot2)
SL.svm.1 = function(..., kernel = "linear") {
  SL.svm(..., kernel = kernel)
}

SL.svm.2 = function(..., kernel = "polynomial", degree = 2, coef0 = 1) {
  SL.svm(..., kernel = kernel, degree = degree, coef0 = coef0)
}

```

```

SL.svm.3 = function(..., kernel = "polynomial", degree = 2, coef0 = 10) {
  SL.svm(..., kernel = kernel, degree = degree, coef0 = coef0)
}

SL.svm.4 = function(..., kernel = "polynomial", degree = 4, coef0 = 1) {
  SL.svm(..., kernel = kernel, degree = degree, coef0 = coef0)
}

SL.svm.5 = function(..., kernel = "polynomial", degree = 4, coef0 = 10) {
  SL.svm(..., kernel = kernel, degree = degree, coef0 = coef0)
}

SL.svm.6 = function(..., kernel = "sigmoid",coef0 = 0) {
  SL.svm(..., kernel = kernel,coef0 = coef0)
}

SL.svm.7 = function(..., kernel = "sigmoid",coef0 = 1) {
  SL.svm(..., kernel = kernel,coef0 = coef0)
}

SL.svm.10 = function(...,type.class = "C-classification", cost = 3) {
  SL.svm(...,type.class = type.class, cost = cost )
}

SL.svm.11 = function(..., kernel = "radial", coef0 = 1) {
  SL.svm(..., kernel = kernel, coef0 = coef0)
}

SL.svm.12 = function(..., kernel = "radial", coef0 = 10) {
  SL.svm(..., kernel = kernel, coef0 = coef0)
}

SL.glmnet.0 <- function(..., alpha = 0,family="binomial"){
  SL.glmnet(..., alpha = alpha , family = family)
}

SL.glmnet.1 <- function(..., alpha = 1,family="binomial"){
  SL.glmnet(..., alpha = alpha , family = family)
}

SL.glmnet.0.25 <- function(..., alpha = 0.25,family="binomial"){
  SL.glmnet(..., alpha = alpha, family = family)
}

SL.glmnet.0.50 <- function(..., alpha = 0.50,family="binomial"){
  SL.glmnet(..., alpha = alpha, family = family)
}

SL.glmnet.0.75 <- function(..., alpha = 0.75,family="binomial"){
  SL.glmnet(..., alpha = alpha, family = family)
}

my_library = c("SL.svm.4",

```

```

"SL.svm.1", "SL.svm.2", "SL.svm.3",
"SL.svm.5", "SL.svm.6", "SL.svm.12",
"SL.svm.7", "SL.glmnet", "SL.svm.10", "SL.svm.11",
"SL.glmnet.0", "SL.glmnet.1", "SL.glmnet.0.50",
"SL.glmnet.0.25", "SL.glmnet.0.75",
"SL.knn", "SL.randomForest", "SL.lm", "SL.mean", "SL.glmnet", "SL.glm", "SL.nnls")

```

The values we are changing are `nu` and `cost(C)`

We are dealing with an almost balanced dataset, so we leave the class weights as default.

Fit `cv.superlearner`

```

sl = SuperLearner(Y = Y_train, X = X_train, family = binomial(),
                  SL.library = my_library)

## Loading required package: glmnet
## Loading required package: Matrix
## Loading required package: foreach
## Loaded glmnet 2.0-13
## Loading required package: class
## Loading required package: randomForest
## randomForest 4.6-12
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:ggplot2':
##
##     margin
## Loading required package: e1071
## Warning in predict.lm(fit, newdata = newX, type = "response"): prediction
## from a rank-deficient fit may be misleading
## Warning: glm.fit: algorithm did not converge
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
## Warning in predict.lm(fit, newdata = newX, type = "response"): prediction
## from a rank-deficient fit may be misleading
## Warning: glm.fit: algorithm did not converge
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

```

```

## Warning in predict.lm(fit, newdata = newX, type = "response"): prediction
## from a rank-deficient fit may be misleading

## Warning: glm.fit: algorithm did not converge

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(fit, newdata = newX, type = "response"): prediction
## from a rank-deficient fit may be misleading

## Warning: glm.fit: algorithm did not converge

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(fit, newdata = newX, type = "response"): prediction
## from a rank-deficient fit may be misleading

## Warning: glm.fit: algorithm did not converge

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(fit, newdata = newX, type = "response"): prediction
## from a rank-deficient fit may be misleading

## Warning: glm.fit: algorithm did not converge

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(fit, newdata = newX, type = "response"): prediction
## from a rank-deficient fit may be misleading

## Warning: glm.fit: algorithm did not converge

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(fit, newdata = newX, type = "response"): prediction
## from a rank-deficient fit may be misleading

## Warning: glm.fit: algorithm did not converge

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(fit, newdata = newX, type = "response"): prediction
## from a rank-deficient fit may be misleading

## Warning: glm.fit: algorithm did not converge

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(fit, newdata = newX, type = "response"): prediction
## from a rank-deficient fit may be misleading

```

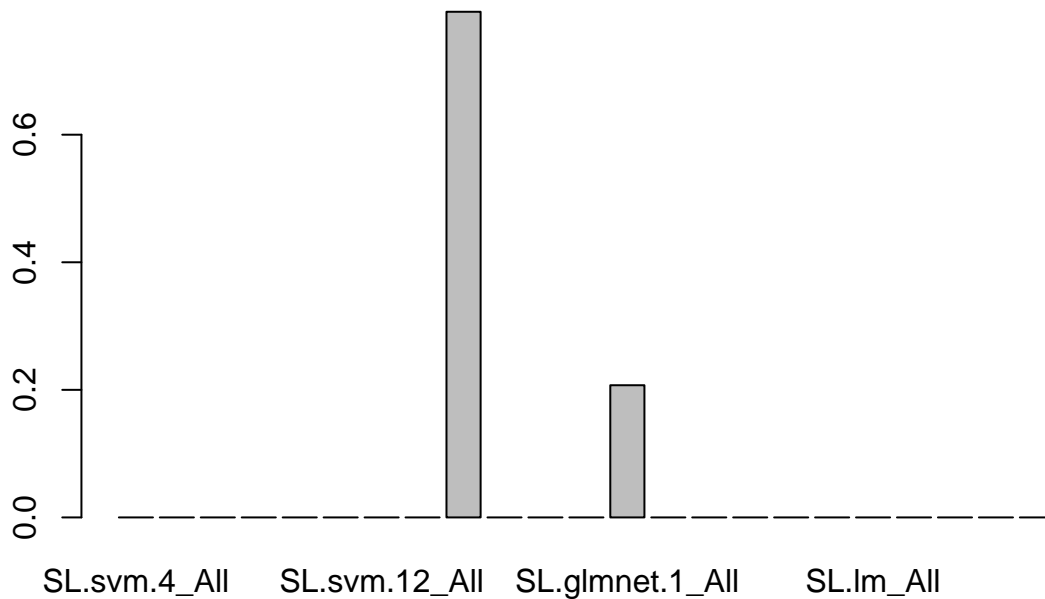
```
## Warning: glm.fit: algorithm did not converge
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
# Review results.
```

```
s1
```

```
##
## Call:
## SuperLearner(Y = Y_train, X = X_train, family = binomial(), SL.library = my_library)
##
##
##
##
```

	Risk	Coef
## SL.svm.4_All	0.2590367	0.0000000
## SL.svm.1_All	0.2573252	0.0000000
## SL.svm.2_All	0.2576456	0.0000000
## SL.svm.3_All	0.2578379	0.0000000
## SL.svm.5_All	0.2569211	0.0000000
## SL.svm.6_All	0.2639106	0.0000000
## SL.svm.12_All	0.2614928	0.0000000
## SL.svm.7_All	0.2612267	0.0000000
## SL.glmnet_All	0.1405737	0.7927469
## SL.svm.10_All	0.2619820	0.0000000
## SL.svm.11_All	0.2556844	0.0000000
## SL.glmnet.0_All	0.2593691	0.0000000
## SL.glmnet.1_All	0.1411372	0.2072531
## SL.glmnet.0.50_All	0.1611375	0.0000000
## SL.glmnet.0.25_All	0.1911410	0.0000000
## SL.glmnet.0.75_All	0.1469679	0.0000000
## SL.knn_All	0.2814714	0.0000000
## SL.randomForest_All	0.2454540	0.0000000
## SL.lm_All	0.5180213	0.0000000
## SL.mean_All	0.2525362	0.0000000
## SL.glmnet_All	0.1416455	0.0000000
## SL.glm_All	0.5267904	0.0000000
## SL.nnls_All	0.4880549	0.0000000

```
barplot(coef(s1))
```



```
pred = predict(s1, X_holdout, onlySL = T)
```

```
# Check the structure of this prediction object.
```

```
str(pred)
```

```
## List of 2
```

```
## $ pred : num [1:74, 1] 0.608 0.453 0.373 0.801 0.899 ...
```

```
## $ library.predict: num [1:74, 1:23] 0 0 0 0 0 0 0 0 0 0 ...
```

```
# We can see which columns are being populated the library.predict.
```

```
summary(pred$library.predict)
```


```
##      V1      V2      V3      V4      V5      V6
## Min.   :0    Min.   :0    Min.   :0    Min.   :0    Min.   :0    Min.   :0
## 1st Qu.:0    1st Qu.:0    1st Qu.:0    1st Qu.:0    1st Qu.:0    1st Qu.:0
## Median :0    Median :0    Median :0    Median :0    Median :0    Median :0
## Mean   :0    Mean   :0    Mean   :0    Mean   :0    Mean   :0    Mean   :0
## 3rd Qu.:0    3rd Qu.:0    3rd Qu.:0    3rd Qu.:0    3rd Qu.:0    3rd Qu.:0
## Max.   :0    Max.   :0    Max.   :0    Max.   :0    Max.   :0    Max.   :0
##      V7      V8      V9      V10     V11
## Min.   :0    Min.   :0    Min.   :0.003555    Min.   :0    Min.   :0
## 1st Qu.:0    1st Qu.:0    1st Qu.:0.166389    1st Qu.:0    1st Qu.:0
## Median :0    Median :0    Median :0.440361    Median :0    Median :0
## Mean   :0    Mean   :0    Mean   :0.458905    Mean   :0    Mean   :0
## 3rd Qu.:0    3rd Qu.:0    3rd Qu.:0.764672    3rd Qu.:0    3rd Qu.:0
## Max.   :0    Max.   :0    Max.   :0.990171    Max.   :0    Max.   :0
##      V12     V13     V14     V15     V16
## Min.   :0    Min.   :0.009035    Min.   :0    Min.   :0    Min.   :0
## 1st Qu.:0    1st Qu.:0.201552    1st Qu.:0    1st Qu.:0    1st Qu.:0
## Median :0    Median :0.448802    Median :0    Median :0    Median :0
## Mean   :0    Mean   :0.462887    Mean   :0    Mean   :0    Mean   :0
## 3rd Qu.:0    3rd Qu.:0.696155    3rd Qu.:0    3rd Qu.:0    3rd Qu.:0
## Max.   :0    Max.   :0.976659    Max.   :0    Max.   :0    Max.   :0
##      V17     V18     V19     V20     V21     V22
## Min.   :0    Min.   :0    Min.   :0    Min.   :0    Min.   :0    Min.   :0
## 1st Qu.:0    1st Qu.:0    1st Qu.:0    1st Qu.:0    1st Qu.:0    1st Qu.:0
```

```
## Median :0   Median :0   Median :0   Median :0   Median :0   Median :0
## Mean    :0   Mean     :0   Mean     :0   Mean     :0   Mean     :0   Mean     :0
## 3rd Qu.:0   3rd Qu.:0   3rd Qu.:0   3rd Qu.:0   3rd Qu.:0   3rd Qu.:0
## Max.    :0   Max.     :0   Max.     :0   Max.     :0   Max.     :0   Max.     :0
##      V23
## Min.     :0
## 1st Qu.:0
## Median   :0
## Mean     :0
## 3rd Qu.:0
## Max.     :0
```

```
# Histogram of our predicted values.
```

```
qplot(pred$pred) + theme_bw()
```


```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```



```
# Scatterplot of original values (0, 1) and predicted values.
```

```
# Ideally we would use jitter or slight transparency to deal with overlap.
```

```
qplot(Y_holdout, pred$pred) + theme_bw()
```

```
# Review AUC - Area Under Curve  
pred_rocr = ROCR::prediction(pred$pred, Y_holdout)  
auc = ROCR::performance(pred_rocr, measure = "auc", x.measure = "cutoff")@y.values[[1]]  
auc
```

```
## [1] 0.8461538
```

The auroc value is 0.86, which is pretty good.